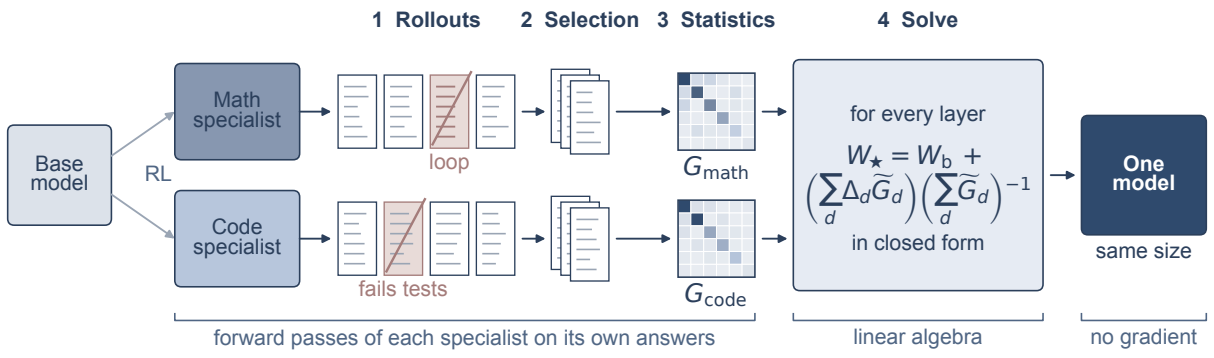
1 Rollouts
2 Selection
3 Statistics
4 Solve
Base model
RL
Math specialist
Code specialist
loop
fails tests
G_{math}
G_{code}
for every layer
$W_\star = W_{\text{b}} + \left(\sum_d \Delta_d \widetilde{G}_d\right)\left(\sum_d \widetilde{G}_d\right)^{-1}$
in closed form
One model
same size
forward passes of each specialist on its own answers
linear algebra
no gradient